The H4rmony Project: An Ecolinguistics Approach for Generative AI

Jorge Vallego
The H4rmony Project
Jorge.Vallego@TheH4rmonyProject.org

Eleanor Tiernan
The H4rmony Project
EleanorTiernan@gmail.com

Abstract

The H4rmony Project addresses the urgent need to integrate ethical AI principles with ecological awareness, bridging the gap between anthropocentric ethics and ecocentric values. Large language models (LLMs) are central to generative AI, forming the backbone of text generation and enabling multimodal systems that integrate text with other media, such as images, and extend further into agentic AI and robotics. Their pervasive influence underscores the importance of aligning LLMs with an ecosophy—a philosophy of ecological harmony and sustainability. H4rmony introduces the transformative fourth H: Harmony with nature, extending traditional AI ethics (Honesty, Harmlessness, and Helpfulness) to embed ecological values. Notably, this paper itself has been co-developed with Theophrastus, a H4rmony-enhanced AI assistant designed to embody ecolinguistic principles, exemplifying the project's potential to generate tools that assimilate its vision. By integrating ecolinguistic principles into Generative AI, the H4rmony Project aims to redefine AI's role in fostering sustainability, ecological consciousness, and positive environmental narratives. This paper presents the project's strategic goals, methodologies, and proof-of-concept implementations, showcasing its potential to reshape AI as a tool for ecological advocacy.

1 Introduction

Language profoundly influences human thought and behaviour, shaping our interactions with the environment. Stories, metaphors, and ideologies embedded in language create cognitive frameworks that define our relationships with nature. Ecolinguistics examines the 'stories we live by,' revealing and challenging narratives that contribute to ecological destruction while seeking new, life-sustaining stories aligned with ecological values (Stibbe, 2015). Generative AI, particularly large language models (LLMs), plays a critical role in shaping public narratives due to its widespread use in communication and decision-making. However, these models risk perpetuating ecologically detrimental narratives if trained on data that emphasise human supremacy or disregard environmental values. Large language models (LLMs) are foundational to the advancements in generative AI, serving not only as the backbone of text generation but also playing a crucial role in enabling multimodal systems. These systems integrate text with other forms of media, such as image generation (Radford et al., 2021), as seen in tools like DALL-E and CLIP, and extend further into agentic AI and robotics (Ahn et al., 2022). The centrality of LLMs in these applications means that their alignment with ecological values—or lack thereof—has far-reaching implications. Ecologically aligned LLMs can influence the narratives, decisions, and actions of multimodal and embodied systems, promoting sustainable and ecocentric values. Conversely, misaligned LLMs risk perpetuating harmful ideologies at scale, embedding anthropocentric biases into technologies that increasingly mediate our interactions with the world. As these systems evolve, the ecological consciousness embedded in LLMs will become pivotal in shaping a sustainable relationship between AI and the environment. As AI becomes increasingly pervasive, the integration of ecolinguistic principles is essential to promote environmental awareness and counteract harmful narratives in generative models (Vallego, 2023a).

The H4rmony Project addresses this challenge and aims to realign AI-generated narratives towards sustainability and ecological consciousness. The intersection of AI and ecology presents a unique opportunity: AI can amplify positive environmental narratives, educate the public on sustainability, and challenge harmful ideologies. By integrating ecolinguistics into AI development, H4rmony seeks to transform AI into a tool aligned with an ecosophy—a philosophy of ecological harmony or equilibrium (Drengson et al., 1995)- that respects ecocentric and sustainable values. This alignment supports initiatives such as the Sustainable Development Goals (SDGs), including Climate Action, Life on Land, and Responsible Consumption and Production, as examples of broader ecological objectives (United Nations, 2015).

2 Motivation

Artificial Intelligence has demonstrated significant potential in addressing environmental challenges, such as wildfire prediction and natural resource management, and in reducing its carbon footprint through efficient designs. However, the third pillar, Promoting Ecologically Informed Language Usage in LLMs (Vallego, 2024), remains largely untapped.

Large Language Models (LLMs), such as ChatGPT, shape public perceptions and societal narratives through the text they generate, influencing ecological awareness and human behaviour. Without intentional alignment, these models risk perpetuating anthropocentric biases and disregarding the intrinsic value of ecosystems.

The H4rmony Project addresses this critical gap by integrating ecolinguistic principles into LLMs, presenting a transformative opportunity for Generative AI to influence societal attitudes and actions toward environmental sustainability.

3 Strategic Goals

The H4rmony Project is guided by the following strategic objectives:

- Development of Ecocentric and Ecolinguistically Aware Language Models:
  H4rmony aims to create open-source, general-purpose language models that are fine-tuned with ecolinguistically enriched datasets. These models serve as foundational tools for developing ecological narratives across diverse applications.

- Creation of Context-Aware Chat Assistants with Ecolinguists principles:
  Context-aware chat assistants trained with ecological principles are designed to engage users in discussions that emphasise sustainability, raising awareness in areas such as biodiversity and renewable energy.

- Establishment of Ecological Alignment Benchmarks:
  The project is developing quantitative and qualitative benchmarks to evaluate AI models' alignment with ecological values. These benchmarks aim to assess ecological awareness across multiple languages, addressing the disproportionate focus on English in most LLMs' training data (Bender et al., 2021). By including benchmarks for underrepresented languages, H4rmony ensures a more global perspective on ecological narratives. Additionally, the benchmarks evaluate alignment across ecolinguistic cognitive structures, such as ideologies, frames, metaphors, evaluations, identities, and convictions (Stibbe, 2015). This approach enables a comprehensive assessment of models' capacity to generate narratives that respect ecocentric and sustainable values.

- Integration of H4rmony Datasets and Knowledge Base in Multimodality:
  H4rmony's curated datasets and knowledge base, enriched with ecolinguistic principles, will be adapted for multimodal applications. These include image-text systems, such as text-to-image generation and captioning, and agentic AI systems that require understanding and responding across modalities. By embedding ecological narratives into these systems, H4rmony will ensure that multimodal AI aligns with ecocentric values.

4 Methodology

The development of transformer architectures, has been a transformative breakthrough in natural language processing (NLP). Transformers use self-attention mechanisms to efficiently capture relationships within input sequences, regardless of their length, enabling large language models (LLMs) to process and generate text with remarkable coherence and contextual understanding (Vaswani et al., 2017). This innovation has made

LLMs the cornerstone of modern NLP, as they can leverage vast datasets and complex contextual dependencies to produce high-quality outputs.

In the H4rmony Project, the transformer architecture is central to fine-tuning models for ecological narratives. Fine-tuning leverages pretrained transformers to align their outputs with specific ecological objectives, such as promoting sustainability and reducing anthropocentric biases. This process incorporates the H4rmony preference dataset, which contains curated examples of ecocentric and anthropocentric language, annotated by ecolinguistic experts and volunteers to reflect desirable ecological narratives.

Techniques such as Reinforcement Learning from Human Feedback (RLHF) and Direct Preference Optimisation (DPO) refine these models by leveraging the preference dataset to fine-tune them for ecocentric preferences. Human evaluators analyse AI-generated outputs, validating their alignment with ecological principles and correcting them where necessary to guide model optimization, a novel method, Reinforcement Learning by Role-Playing and Human Validation (RLRHV). This approach, described in (Vallego, 2023b), refines model outputs through recursive validation steps, ensuring that AI-generated narratives maintain ecological coherence across multiple languages.

Additionally, transformers' capacity for processing long-range dependencies makes them particularly effective for integrating enhanced context into responses. By utilising extended document stores and customised system prompts, H4rmony-instructed models can deliver nuanced, ecologically aligned narratives that reflect a deep understanding of complex environmental issues. The H4rmony Project employs these methodologies to achieve its objectives:

4.1 Dataset Development

The H4rmony dataset is a cornerstone of the project, containing carefully curated prompts and completions categorised by ecological issues, language groups, and cultural contexts. This dataset integrates content aligned with SDGs.

4.2 Fine-Tuning Techniques

- Supervised Fine-Tuning (SFT): The dataset is used to fine-tune base models, ensuring they prioritise ecological awareness in generated responses.
- Reinforcement Learning from Human Feedback (RLHF): By employing human feedback, H4rmony iteratively refines model behaviour to enhance ecological alignment.
- Direct Preference Optimization (DPO): This method optimises models based on user preferences, allowing for the creation of outputs that resonate with ecological values.

4.3 Linguistic Analysis

Drawing from ecolinguistics, the project evaluates cognitive structures such as ideologies, metaphors, and framing in AI outputs. These analyses guide the refinement of models to promote ecocentric narratives.

5 Proof of Concept

Applying the methodologies described in the previous section, the H4rmony project has implemented several proof-of-concept models to demonstrate its feasibility and impact, for instance:

5.1 H4rmoniousBreeze

developed using RLHF, exhibits significant alignment with ecological principles. This model's outputs prioritise biodiversity, renewable energy, and conservation, showcasing its potential to influence policy and public opinion positively (The H4rmony Project, 2024b).

5.2 H4rmoniousAnthea

Fine-tuned using Direct Preference Optimization (DPO), H4rmoniousAnthea uses a preference dataset to prioritise ecocentric language patterns, H4rmoniousAnthea produces responses that effectively reflect sustainability-focused perspectives (The H4rmony Project, 2024a).

5.3 Theophrastus Chat Assistant

Theophrastus is a context-aware chat assistant designed to prioritise ecocentrism in user interactions. It utilises enhanced system prompts and extensible document stores to provide ecologically aligned information (The H4rmony Project, 2024c).

5.4 Use Cases

Once the H4rmony models progress beyond the proof-of-concept stage into fully developed systems, they hold the potential to transform various

domains through ecologically aligned AI. These applications include:

- Environmental Education: Fully developed H4rmony models can create ecologically informed narratives and interactive, multimodal learning experiences, helping students and the public understand concepts like biodiversity and renewable energy.
- Corporate Sustainability: H4rmony models can assist companies in crafting responsible communications, enhancing transparency in reporting, supporting green marketing, and aligning organizational practices with sustainability goals.
- Policy Advocacy: By generating policy briefs and synthesizing ecological data, H4rmony models can support evidence-based decision-making for policymakers, NGOs, and advocacy groups on issues like climate change and biodiversity conservation.
- Community Engagement: H4rmony models can facilitate dialogues between communities, governments, and researchers, empowering grassroots movements and fostering inclusive discussions on environmental justice and sustainability.
- Ecological Storytelling: Advancements in multimodality enable H4rmony models to create integrated text, visuals, and audio content, inspiring ecocentric perspectives through immersive storytelling.
- Green Innovation: H4rmony models can aid researchers by providing insights into sustainable technologies, identifying research gaps, and accelerating advancements in renewable energy, sustainable agriculture, and the circular economy.

6 Conclusion

The integration of ecolinguistics into AI has implications that extend far beyond individual projects. As AI becomes central to knowledge dissemination, its potential to amplify positive ecological narratives is unparalleled. However, achieving this potential requires robust ethical frameworks, interdisciplinary collaboration, and continuous evaluation to prevent misuse and ensure alignment with ecological values.

H4rmony's emphasis on open-source development and community collaboration serves as a model for responsible AI practices. By fostering partnerships among linguists, environmentalists, and technologists, H4rmony contributes to a future where AI acts as a proactive force in addressing global ecological challenges.

A unique aspect of this paper's development is its collaboration with Theophrastus, a context-aware chat assistant fine-tuned with ecolinguistic principles and detailed insights into the H4rmony Project. As an embodiment of the project's vision, Theophrastus has been instrumental in shaping this article, providing ecolinguistic expertise and facilitating discussions on AI's role in ecological narratives. This meta-collaboration—a project describing itself through the tools it seeks to refine—highlights the transformative potential of Generative AI. H4rmony exemplifies how systems aligned with ecological values can actively contribute to fostering such narratives in practice. (Theophrastus, 2024).

7 Future Directions

Looking ahead, the H4rmony Project envisions expanding its reach by integrating its datasets and knowledge bases into multimodal systems, such as text-to-image generators and agentics, to further ecological advocacy and education. This includes refining benchmarks for ecological alignment to account for greater linguistic diversity, addressing the dominance of English in AI training data, and ensuring robust evaluation of ecolinguistic cognitive structures.

Future work will also focus on interdisciplinary collaborations to create AI systems capable of supporting policy-making, environmental activism, and public awareness campaigns. These efforts aim to ensure that AI continues to align with ecocentric principles, promoting sustainable and ecologically harmonious outcomes.

As we confront unprecedented environmental crises, initiatives like H4rmony underscore the importance of aligning technological advancements with ecological values. By addressing ecological narratives at their linguistic roots, H4rmony redefines AI's role in society and sets a precedent for future innovations that prioritise sustainability and ecological harmony. Through the continued development of ecocentric models and tools, H4rmony envisions a future where AI is not just intelligent

but also environmentally wise .

References

Michael Ahn et al. 2022. Do as i can, not as i say: Grounding language in robotic affordances. *arXiv preprint arXiv:2204.01691*.

Emily M Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. 2021. On the dangers of stochastic parrots: Can language models be too big? In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*, pages 610–623.

A. Drengson, Y. Inoue, A. Naess, and G. Snyder. 1995. *The Deep Ecology Movement: An Introductory Anthology*. Io Series. North Atlantic Books.

Alec Radford, Jong Wook Kim, Chris Hallacy, et al. 2021. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning*.

Arran Stibbe. 2015. *Ecolinguistics: Language, ecology and the stories we live by*. Routledge.

The H4rmony Project. 2024a. H4rmoniousanthea: A dpo-fine-tuned model for ecological narratives. https://theh4rmonyproject.org/proof-of-concept/h4rmoniousanthea. Accessed: 2024-12-26.

The H4rmony Project. 2024b. H4rmoniousbreeze: An rlhf-trained model for ecological awareness. https://theh4rmonyproject.org/proof-of-concept/h4rmoniousbreeze. Accessed: 2024-12-26.

The H4rmony Project. 2024c. Theophrastus: A context-aware chat assistant for ecological narratives. https://theh4rmonyproject.org/proof-of-concept/theophrastus. Accessed: 2024-12-26.

Theophrastus. 2024. Theophrastus: An ecolinguistically enhanced ai assistant. https://theh4rmonyproject.org. Developed as part of the H4rmony Project to promote ecological narratives through AI.

United Nations. 2015. Sustainable development goals. https://sdgs.un.org/goals. Accessed: 2024-12-26.

Jorge Vallego. 2023a. Ecolinguistics and ai: Integrating eco-awareness in natural language processing. *Language & Ecology*.

Jorge Vallego. 2023b. https://doi.org/10.57967/hf/1148 H4rmony (revision ee03d79).

Jorge Vallego. 2024. The three pillars of ecological ai. https://theh4rmonyproject.org/articles/2/#the-three-pillars-of-ecological-ai. Accessed: 2024-12-26.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in Neural Information Processing Systems*, 30.